

Education

Georgia Institute of Technology Atlanta, USA Ph.D. (CS, minor in Optimization) 2013
Indian Institute of Technology Kharagpur, India M.Sc. & B.Sc. (Math & Computing) 2007

Research & Professional Experience

2018 – present Principal Research Staff Member, IBM Research
2017 – 2018 Principal Product Architect, Infosys Ltd
2015 – 2017 Senior Staff Research Scientist, Skytree Inc
2013 – 2015 Member of Technical Staff, Skytree Inc

Research Interests

I have a wide range of research interests, and I am always excited to learn about new areas and make connections between different areas. In this process, I have spent some time focusing on the following areas: Optimization (single and bilevel), Automated Machine Learning & Data Science, Large Scale Learning, Computational Geometry, Efficient All-Pairs Algorithms & Analysis, Density Estimation, Kernel Methods, Associative Memories & Energy-based Models, Machine Unlearning, Sparse Learning, Neuro-inspired Learning, Compositional Generalization.

Selected Publications

I have been fortunate to publish in various research areas at top AI, Machine Learning and Data Mining conferences (h-index: 23). I have also filed 20+ patents. Here is a selected subset:

- B. Hoover, D. H. Chau, H. Strobelt, P. Ram, and D. Krotov. Dense associative memory through the lens of random features. *Advances in Neural Information Processing Systems*, 37, 2024
- J. Jia, J. Liu, Y. Zhang, P. Ram, N. Baracaldo, and S. Liu. Wagle: Strategic weight attribution for effective and modular unlearning in large language models. *Advances in Neural Information Processing Systems*, 37, 2024
- P. Ram, T. Klinger, and A. G. Gray. What makes models compositional? a theoretical view. In *Proceedings of the Thirty-First International Joint Conference on Artificial Intelligence (IJCAI-24)*, 2024
- J. Liu, P. Ram, Y. Yao, G. Liu, Y. Liu, P. Sharma, S. Liu, et al. Model sparsity can simplify machine unlearning. *Advances in Neural Information Processing Systems*, 36, 2023 spotlight, 90+ citations
- B. Saha, D. Krotov, M. J. Zaki, and P. Ram. End-to-end differentiable clustering with associative memories. In *International Conference on Machine Learning*, pages 29649–29670. PMLR, 2023
- Y. Zhou, P. Ram, T. Salonidis, N. Baracaldo, H. Samulowitz, and H. Ludwig. Single-shot general hyper-parameter optimization for federated learning. In *The 11th International Conference on Learning Representations*, 2023 notable top-25%
- A. Gu, S. Lu, P. Ram, and T.-W. Weng. Min-max multi-objective bilevel optimization with applications in robust machine learning. In *The 11th International Conference on Learning Representations*, 2023
- P. Ram and K. Sinha. Federated nearest neighbor classification with a colony of fruit-flies. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pages 8036–8044, 2022
- K. Sinha and P. Ram. Fruit-fly inspired neighborhood encoding for classification. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*, pages 1470–1480, 2021
- S. Liu, P. Ram, D. Vijaykeerthy, D. Bouneffouf, G. Bramble, H. Samulowitz, D. Wang, A. Conn, and A. Gray. An admm based framework for automl pipeline configuration. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 4892–4899, 2020 90 citations
- P. Ram and K. Sinha. Revisiting kd-tree for nearest neighbor search. In *Proceedings of the 25th acm sigkdd international conference on knowledge discovery & data mining*, pages 1378–1388, 2019 140 citations
- R. R. Curtin, P. Ram, and A. G. Gray. Fast exact max-kernel search. In *Proceedings of the 2013 SIAM International Conference on Data Mining*, pages 1–9. Society for Industrial and Applied Mathematics, 2013 best paper finalist
- P. Ram and A. G. Gray. Maximum inner-product search using cone trees. In *Proceedings of the 18th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 931–939, 2012 200+ citations
- P. Ram and A. G. Gray. Density estimation trees. In *Proceedings of the 17th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 627–635, 2011 140 citations
- P. Ram, D. Lee, W. March, and A. Gray. Linear-time algorithms for pairwise statistical problems. *Advances in Neural Information Processing Systems*, 22, 2009 spotlight, 90 citations

Recognition

- IBM Research Accomplishments: *Towards Automating AI Lifecycle with AutoAI* (2019), *AutoAI: The Next Level* (2020), *Federated Learning Security & Privacy* (2022), *Neurobiologically Inspired Foundation Models* (2023).
- TMLR Expert Reviewer (2023), NeurIPS Top Reviewer (2022, 2023), ICML Top Reviewer (2022).
- IBM Invention Plateau 1 (May 2020) - 5 (January 2024).
- IBM Research 2020 Pat Goldberg Best Paper Finalist for *An ADMM based Framework for AutoML Pipeline Configuration*.
- SIAM Data Mining 2013 Best Paper Finalist for *Fast Exact Max-Kernel Search*.

Service

- Organizer, NeurIPS 2023 Workshop on [Associative Memory and Hopfield Networks](#).
- Senior PC member AAAI (2024, 2025), NeurIPS Area Chair (2024), ICML Area Chair (2025).
- Office-hours, ICLR 2023.
- Organizer, KDD 2022 Tutorial on [Gradual AutoML using Lale](#).
- Regular PC member for NeurIPS, ICML, AISTATS, ICLR, UAI, KDD, AAAI conferences.

Invited Talks

- ➡ SOON Dense Associative Memories through the Lens of Random Features, **VSAONLINE**, **June 2025**.
- ➡ SOON Associative Memories and its Role in Machine Learning, **ICLR 2025** workshop on *New Frontiers in Associative Memories*.
- The Critical Role of Fundamental Mathematical Tools on the Path to AGI, **IIITDM Kancheepuram** 2024 workshop on *Future Perspectives of AI and Data Sciences : Algorithms and Applications*. [\[video\]](#)
- What makes Models Compositional? A Neuro-Symbolic Theoretical View, **IJCAI 2024** workshop on *Logical Foundations of Neuro-Symbolic AI*.
- Robust Multi-objective Bilevel Optimization With Applications In Machine Learning, **INFORMS 2022** session on *Bilevel Stochastic Methods for Optimization and Learning*.
- Density Estimation Trees, **SIAM Uncertainty Quantification 2014 (UQ14)** mini-symposium on *Non-parametric Density Estimation*.
- Max-kernel Search: How to search for just about anything?, **MLConf Atlanta 2014**. [\[video\]](#)

Teaching

- Bilevel Optimization at the *Polyhedra and Combinatorial Optimization Days 2023 (JPOC13) summer school on Combinatorial Optimization & Machine Learning*.
- Compositional Generalization at the *Neuro-Symbolic Summer School 2023*.
- Hands-on tutorial at the *KDD 2022 Tutorial on Gradual AutoML using Lale*.
- Teaching Assistant, *Data and Visual Analytics* (Spring 2011, 2013), Georgia Tech.

Research Grants

- *Fast Inference and Alignment for Large Multi-modal Models*, RPI-IBM Future of Computing Research Collaboration Program, 2025
- *Meta-Transfer-Learning for Tabular Data Distillation, Generation, and Predictive Modeling*, RPI-IBM Future of Computing Research Collaboration Program, 2025
- *FIT: Fast Inference using Transformer Models*, RPI-IBM Future of Computing Research Collaboration Program, 2024
- *Data Distillation in Tabular Data: A Foundation Model Approach*, RPI-IBM Future of Computing Research Collaboration Program, 2024

-
- *A Framework for Automating Decentralized Training of Foundation Models*, RPI-IBM AI Research Collaboration Program, 2023
 - *AutoDML: A Framework for Automating Decentralized Machine Learning*, RPI-IBM AI Research Collaboration Program, 2022

Mentorship

I have mentored students both during summer internships at IBM research, and through continued collaborations with their advisors at various academic institutions.

Continuing projects

- Ben Hoover (GT), 2024-current.
- Inwon Kang (RPI), 2023-current.
- Jinghan Jia (MSU), 2023-present.
- Bishwajit Saha (RPI), 2022-2024.
- Momin Abbas (RPI), 2022-2023.
- Jiancheng Liu (MSU), 2022-current.
- Yuguang Yao (MSU), 2020-present.
- Yihua Zhang (MSU), 2020-present.
- Yunfei Teng (NYU), 2020-2022.

Summer interns at IBM

- Inwon Kang (RPI) 2023, 2024
- Charlotte Park (MIT) 2024
- Momin Abbas (RPI) 2023
- Bishwajit Saha (RPI) 2022, 2023
- Xinying Qi (RPI) 2022, 2023
- Lucky Yerimah (RPI) 2021

Undergraduate researchers at MIT-IBM

- Alex Gu (MIT) 2021

Publications

Here is the complete list of my publications in chronological order for each publication type.

Book Chapters

1. Y. Zhou, P. Ram, T. Salonidis, N. Baracaldo, H. Samulowitz, and H. Ludwig. Hyper-parameter optimization in federated learning. In *Federated Learning*, pages 237–255. Academic Press, 2024

Journals

1. R. R. Curtin, J. R. Cline, N. P. Slagle, W. B. March, P. Ram, N. A. Mehta, and A. G. Gray. Mlpack: A scalable c++ machine learning library. *The Journal of Machine Learning Research*, 14(1):801–805, 2013
2. R. R. Curtin and P. Ram. Dual-tree fast exact max-kernel search. *Statistical Analysis and Data Mining: The ASA Data Science Journal*, 7(4):229–253, 2014
3. R. R. Curtin, D. Lee, W. B. March, and P. Ram. Plug-and-play dual-tree algorithm runtime analysis. *Journal of Machine Learning Research*, 16:3269–3297, 2015
4. O. Keivani, K. Sinha, and P. Ram. Improved maximum inner product search with better theoretical guarantee using randomized partition trees. *Machine Learning*, 107(6):1069–1094, 2018

Conferences

1. P. Ram, D. Lee, W. March, and A. Gray. Linear-time algorithms for pairwise statistical problems. *Advances in Neural Information Processing Systems*, 22, 2009
2. P. Ram, D. Lee, H. Ouyang, and A. Gray. Rank-approximate nearest neighbor search: Retaining meaning and speed in high dimensions. *Advances in Neural Information Processing Systems*, 22, 2009

3. W. B. March, P. Ram, and A. G. Gray. Fast euclidean minimum spanning tree: algorithm, analysis, and applications. In *Proceedings of the 16th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 603–612, 2010
4. P. Ram, D. Lee, and A. G. Gray. Nearest-neighbor search on a time budget via max-margin trees. In *SIAM Data Mining*, 2012
5. P. Ram and A. G. Gray. Density estimation trees. In *Proceedings of the 17th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 627–635, 2011
6. P. Ram and A. G. Gray. Maximum inner-product search using cone trees. In *Proceedings of the 18th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 931–939, 2012
7. N. Koenigstein, P. Ram, and Y. Shavitt. Efficient retrieval of recommendations in a matrix factorization framework. In *Proceedings of the 21st ACM international conference on Information and knowledge management*, pages 535–544, 2012
8. R. R. Curtin, P. Ram, and A. G. Gray. Fast exact max-kernel search. In *Proceedings of the 2013 SIAM International Conference on Data Mining*, pages 1–9. Society for Industrial and Applied Mathematics, 2013
9. R. Curtin, W. March, P. Ram, D. Anderson, A. Gray, and C. Isbell. Tree-independent dual-tree algorithms. In *International Conference on Machine Learning*, pages 1435–1443. PMLR, 2013
10. P. Ram and A. Gray. Which space partitioning tree to use for search? *Advances in Neural Information Processing Systems*, 26, 2013
11. O. Keivani, K. Sinha, and P. Ram. Improved maximum inner product search with better theoretical guarantees. In *2017 International Joint Conference on Neural Networks (IJCNN)*, pages 2927–2934. IEEE, 2017
12. P. Ram and K. Sinha. Revisiting kd-tree for nearest neighbor search. In *Proceedings of the 25th acm sigkdd international conference on knowledge discovery & data mining*, pages 1378–1388, 2019
13. D. Wang, J. D. Weisz, M. Muller, P. Ram, W. Geyer, C. Dugan, Y. Tausczik, H. Samulowitz, and A. Gray. Human-ai collaboration in data science: Exploring data scientists’ perceptions of automated ai. *Proceedings of the ACM on human-computer interaction*, 3(CSCW):1–24, 2019
14. S. Liu, P. Ram, D. Vijaykeerthy, D. Bouneffouf, G. Bramble, H. Samulowitz, D. Wang, A. Conn, and A. Gray. An admm based framework for automl pipeline configuration. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 4892–4899, 2020
15. D. Wang, P. Ram, D. K. I. Weidele, S. Liu, M. Muller, J. D. Weisz, A. Valente, A. Chaudhary, D. Torres, H. Samulowitz, et al. Autoai: Automating the end-to-end ai lifecycle with humans-in-the-loop. In *Companion Proceedings of the 25th International Conference on Intelligent User Interfaces*, pages 77–78, 2020
16. M. Katz, P. Ram, S. Sohrabi, and O. Udrea. Exploring context-free languages via planning: The case for automating machine learning. In *Proceedings of the International Conference on Automated Planning and Scheduling*, volume 30, pages 403–411, 2020
17. D. Bouneffouf, C. Aggarwal, T. Hoang, U. Khurana, H. Samulowitz, B. Buesser, S. Liu, T. Pedapati, P. Ram, A. Rawat, et al. Survey on automated end-to-end data science? In *2020 International Joint Conference on Neural Networks (IJCNN)*, pages 1–9. IEEE, 2020
18. R. Marinescu, A. Kishimoto, P. Ram, A. Rawat, M. Wistuba, P. P. Palmes, and A. Botea. Searching for machine learning pipelines using a context-free grammar. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 8902–8911, 2021
19. K. Sinha and P. Ram. Fruit-fly inspired neighborhood encoding for classification. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*, pages 1470–1480, 2021
20. A. Kishimoto, D. Bouneffouf, R. Marinescu, P. Ram, A. Rawat, M. Wistuba, P. Palmes, and A. Botea. Bandit limited discrepancy search and application to machine learning pipeline optimization. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pages 10228–10237, 2022
21. G. Baudart, M. Hirzel, K. Kate, P. Ram, A. Shinnar, and J. Tsay. Pipeline combinators for gradual automl. *Advances in Neural Information Processing Systems*, 34:19705–19718, 2021
22. A. Kannan, A. R. Choudhury, V. Saxena, S. Raje, P. Ram, A. Verma, and Y. Sabharwal. Hyperaspo: Fusion of model and hyper parameter optimization for multi-objective machine learning. In *2021 IEEE International Conference on Big Data (Big Data)*, pages 790–800. IEEE, 2021
23. P. Ram and K. Sinha. Federated nearest neighbor classification with a colony of fruit-flies. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 36, pages 8036–8044, 2022
24. P. Ram. On the optimality gap of warm-started hyperparameter optimization. In *International Conference on Automated Machine Learning*, pages 12–1. PMLR, 2022
25. D. Subramanian, S. Wasserkrug, P. Murali, D. Phan, P. Ram, O. Davidovich, X. Ceugniet, and F. Katai. Data and knowledge driven optimization model generation for flow based optimization problems. In *INFORMS Annual Meeting*, 2021
26. P. Zhao, P. Ram, S. Lu, Y. Yao, D. Bouneffouf, X. Lin, and S. Liu. Learning to generate image source-agnostic universal adversarial perturbations. In *Proceedings of the Thirty-First International Joint Conference on Artificial*

Intelligence (IJCAI-22), 2022

27. M. Hirzel, K. Kate, P. Ram, A. Shinnar, and J. Tsay. Gradual automl using lale. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 4794–4795, 2022
28. Y. Teng, A. Choromanska, M. Campbell, S. Lu, P. Ram, and L. Horesh. Overcoming catastrophic forgetting via direction-constrained optimization. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, pages 675–692. Springer International Publishing Cham, 2022
29. A. Gu, S. Lu, P. Ram, and L. Weng. Robust multi-objective bilevel optimization with applications in machine learning. In *INFORMS Annual Meeting*, 2022
30. Y. Zhang, Y. Yao, P. Ram, P. Zhao, T. Chen, M. Hong, Y. Wang, and S. Liu. Advancing model pruning via bi-level optimization. *Advances in Neural Information Processing Systems*, 35:18309–18326, 2022
31. P. Palmes, A. Kishimoto, R. Marinescu, P. Ram, and E. Daly. Distributed automl pipeline search in pc/raspi k8s cluster. In *Julia Con*, 2022
32. P. Ram, A. G. Gray, H. C. Samulowitz, and G. Bramble. Toward theoretical guidance for two common questions in practical cross-validation based hyperparameter selection. In *Proceedings of the 2023 SIAM International Conference on Data Mining (SDM)*, pages 802–810. Society for Industrial and Applied Mathematics, 2023
33. Y. Zhang, P. Sharma, P. Ram, M. Hong, K. Varshney, and S. Liu. What is missing in irm training and evaluation? challenges and solutions. In *The 11th International Conference on Learning Representations*, 2023
34. A. Gu, S. Lu, P. Ram, and T.-W. Weng. Min-max multi-objective bilevel optimization with applications in robust machine learning. In *The 11th International Conference on Learning Representations*, 2023
35. Y. Zhou, P. Ram, T. Salonidis, N. Baracaldo, H. Samulowitz, and H. Ludwig. Single-shot general hyperparameter optimization for federated learning. In *The 11th International Conference on Learning Representations*, 2023
36. P. Dube, T. Salonidis, P. Ram, and A. Verma. Runtime prediction of machine learning algorithms in automl systems. In *ICASSP 2023-2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 1–5. IEEE, 2023
37. B. Saha, D. Krotov, M. J. Zaki, and P. Ram. End-to-end differentiable clustering with associative memories. In *International Conference on Machine Learning*, pages 29649–29670. PMLR, 2023
38. M. Feffer, M. Hirzel, S. C. Hoffman, K. Kate, P. Ram, and A. Shinnar. Searching for fairer machine learning ensembles. In *International Conference on Automated Machine Learning*, pages 17–1. PMLR, 2023
39. M. Lazuka, A. Anghel, P. Ram, H. Pozidis, and T. Parnell. xcloudserving: Automated ml serving across clouds. In *2023 IEEE 16th International Conference on Cloud Computing (CLOUD)*, pages 1–12. IEEE, 2023
40. J. Liu, P. Ram, Y. Yao, G. Liu, Y. Liu, P. Sharma, S. Liu, et al. Model sparsity can simplify machine unlearning. *Advances in Neural Information Processing Systems*, 36, 2023
41. I. Kang, P. Ram, Y. Zhou, H. Samulowitz, and O. Seneviratne. Effective data distillation for tabular datasets (student abstract). In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 23533–23534, 2024
42. M. Abbas, Y. Zhou, P. Ram, N. Baracaldo, H. Samulowitz, T. Salonidis, and T. Chen. Enhancing in-context learning via linear probe calibration. In *International Conference on Artificial Intelligence and Statistics*, pages 307–315. PMLR, 2024
43. M. Abbas, Y. Zhou, N. Baracaldo, H. Samulowitz, P. Ram, and T. Salonidis. Byzantine-resilient bilevel federated learning. In *2024 IEEE 13rd Sensor Array and Multichannel Signal Processing Workshop (SAM)*, pages 1–5. IEEE, 2024
44. P. Ram, T. Klinger, and A. G. Gray. What makes models compositional? a theoretical view. In *Proceedings of the Thirty-First International Joint Conference on Artificial Intelligence (IJCAI-24)*, 2024
45. J. Jia, J. Liu, Y. Zhang, P. Ram, N. Baracaldo, and S. Liu. Wagle: Strategic weight attribution for effective and modular unlearning in large language models. *Advances in Neural Information Processing Systems*, 37, 2024
46. B. Hoover, D. H. Chau, H. Strobelt, P. Ram, and D. Krotov. Dense associative memory through the lens of random features. *Advances in Neural Information Processing Systems*, 37, 2024
47. S. Carrow, K. H. Erwin, O. Vilenskaia, P. Ram, T. Klinger, N. A. Khan, N. Makondo, and A. Gray. Neural reasoning networks: Efficient interpretable neural networks with automatic textual explanations. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2025

Workshops

1. P. Ram and A. G. Gray. Fraud detection with density estimation trees. In *KDD 2017 Workshop on Anomaly Detection in Finance*, pages 85–94. PMLR, 2018
2. P. Ram, S. Liu, D. Vijaykeerthi, D. Wang, D. Bouneffouf, G. Bramble, H. Samulowitz, and A. G. Gray. Solving constrained cash problems with admm. In *7th ICML 2020 Workshop on Automated Machine Learning (AutoML)*, 2020

3. G. Baudart, M. Hirzel, K. Kate, P. Ram, and A. Shinnar. Lale: Consistent automated machine learning. In *AutoML Workshop @ KDD'20*, 2020
4. M. Hirzel, K. Kate, and P. Ram. Engineering fair machine learning pipelines. In *ICLR 2021 Workshop on Responsible AI*, 2021
5. P. Ram, A. G. Gray, and H. Samulowitz. Leveraging theoretical tradeoffs in hyperparameter selection for improved empirical performance. In *8th ICML 2021 Workshop on Automated Machine Learning (AutoML)*, 2021
6. C. Fan, P. Ram, and S. Liu. Sign-MAML: Efficient model-agnostic meta-learning by SignSGD. In *5th NeurIPS 2021 Workshop on Meta-Learning*, 2021
7. S. Wasserkrug, O. Davidovich, E. Shindin, D. Subramanian, P. Ram, P. Murali, D. Phan, N. Zhou, and L. M. Nguyen. Ensuring the quality of optimization solutions in data generated optimization models. In *IJCAI 2021 DSO (Data Science Meets Optimisation) Workshop*, 2021
8. P. Ram and K. Sinha. Flynn: Fruit-fly inspired federated nearest neighbor classification. In *International Workshop on Federated Learning for User Privacy and Data Confidentiality in Conjunction with ICML 2021 (FL-ICML'21)*, 2021
9. Y. Zhou, P. Ram, T. Salonidis, N. Baracaldo, H. Samulowitz, and H. Ludwig. Flora: Single-shot hyper-parameter optimization for federated learning. In *NeurIPS 2021 Workshop on New Frontiers in Federated Learning: Privacy, Fairness, Robustness, Personalization and Data Ownership*, 2021
10. A. Gu, S. Lu, P. Ram, and L. Weng. Nonconvex min-max bilevel optimization for task robust meta learning. In *ICML 2021 Workshop on Beyond first-order methods in ML systems*, 2021
11. M. Feffer, M. Hirzel, S. C. Hoffman, K. Kate, P. Ram, and A. Shinnar. An empirical study of modular bias mitigators and ensembles. In *ICML 2022 Workshop on Benchmarking Data for Data-Centric AI (DataPerf@ICML)*, 2022
12. R. Marinescu, T. Pedapati, L. Vu, P. Palmes, T. Mummert, P. Kirchner, D. Subramanian, P. Ram, and D. Boun-effouf. Automated decision optimization with reinforcement learning. In *AAAI-22 Workshop on AI for Decision Optimization*, 2022
13. O. Davidovich, P. Ram, S. Wasserkrug, S. Subramaniam, N. Zhou, D. Phan, P. Murali, and L. Nguyen. Addressing solution quality in data generated optimization models. In *AAAI-22 Workshop on AI for Decision Optimization*, 2022
14. T. Klinger, L. Liu, S. Dan, M. Crouse, P. Ram, and A. Gray. Compositional program generation for systematic generalization. In *IJCAI 2023 Workshop on Knowledge-Based Compositional Generalization*, 2023
15. P. Ram, T. Klinger, and A. G. Gray. How compositional is a model? In *IJCAI 2023 Workshop on Knowledge-Based Compositional Generalization*, 2023
16. M. Hirzel and P. Ram. Oversampling to repair bias and imbalance simultaneously. In *Workshop track at AutoML Conference 2023*, 2023
17. P. Ram, T. Klinger, and A. Gray. What makes models compositional? a neuro-symbolic theoretical view. In *First IJCAI 2024 International Workshop on Logical Foundations of Neuro-Symbolic AI*, 2024
18. B. Saha, D. Krotov, M. J. Zaki, and P. Ram. Deep clustering with associative memories. In *NeurIPS 2024 Workshop on Machine Learning and Compression*, 2024

Patents

1. M. Gibiansky, R. Riegel, Y. Yang, P. Ram, and A. Gray. Configurable machine learning method selection and parameter optimization system and method, Apr. 21 2016. US Patent App. 14/883,522
2. D. K. I. Weidele and P. Ram. Conditional parallel coordinates, July 27 2021. US Patent 11,074,728
3. P. Ram, D. Wang, D. Vijaykeerthy, V. Saxena, S. Liu, A. Chaudhary, G. Bramble, H. C. Samulowitz, and A. Gray. Machine learning with multiple constraints, Mar. 10 2022. US Patent App. 17/015,243
4. M. Katz, P. Ram, S. S. Araghi, and O. Udrea. Machine learning pipeline generation, June 16 2022. US Patent App. 17/119,134
5. A. Kannan, V. Saxena, A. R. Choudhury, Y. Sabharwal, P. Ram, A. Verma, and S. M. Raje. Multi-objective machine learning with model and hyperparameter optimization fusion, Mar. 9 2023. US Patent App. 17/470,763
6. B. Zhang, G. Bramble, P. Ram, and H. C. Samulowitz. Methods for automatically configuring performance evaluation schemes for machine learning algorithms, June 20 2023. US Patent 11,681,931
7. D. K. I. Weidele, P. Ram, D. Wang, A. N. Valente, and A. Chaudhary. Conditional parallel coordinates in automated artificial intelligence with constraints, Jan. 17 2023. US Patent 11,556,816
8. A. Kishimoto, D. Bouneffouf, B. Chen, R. Marinescu, P. Ram, A. Rwat, and M. Wistuba. Automated generation of a machine learning pipeline, Apr. 11 2023. US Patent 11,625,632
9. L. Horesh, G. Nannicini, O. Gunluk, S. Dash, P. Ram, and A. Gray. Free-form integration of machine learning model primitives, Mar. 7 2023. US Patent 11,599,829
10. R. Marinescu and P. Ram. Automl with multiple objectives and tradeoffs thereof, Mar. 30 2023. US Patent App. 17/489,781
11. R. Marinescu and P. Ram. Enhanced machine learning pipelines with multiple objectives and tradeoffs, Apr. 13

-
2023. US Patent App. 17/450,670
 12. Y. Zhou, P. Ram, N. B. Angel, T. Salonidis, H. C. Samulowitz, M. Wistuba, and H. H. Ludwig. Performing automated tuning of hyperparameters in a federated learning environment, June 15 2023. US Patent App. 17/547,122
 13. E. S. Wasserkrug, O. Davidovich, E. Shindin, D. Subramanian, and P. Ram. Augmenting mathematical optimization models generated from historical data, Aug. 3 2023. US Patent App. 17/589,092
 14. R. Marinescu, P. Ram, D. Bouneffouf, T. Pedapati, and P. Palmes. Computing robust policies in offline reinforcement learning, Dec. 21 2023. US Patent App. 17/807,397
 15. V. Saxena, A. Kannan, S. M. Raje, P. Ram, Y. Sabharwal, and A. Verma. Multi-objective automated machine learning, Oct. 8 2024. US Patent 12,112,249
 16. R. Marinescu, P. Ram, D. Bouneffouf, T. Pedapati, and P. Palmes. Reinforcement learning with multiple objectives and tradeoffs, Apr. 25 2024. US Patent App. 17/972,291
 17. Y. J. Ong, Y. Zhou, P. Ram, T. Salonidis, and N. B. Angel. Personalized federated learning of gradient boosted trees, May 2 2024. US Patent App. 18/175,006
 18. Y. Zhou, P. Ram, T. Salonidis, N. B. Angel, H. C. Samulowitz, and H. H. Ludwig. Automated tuning of hyperparameters based on rankings in a federated learning environment, May 2 2024. US Patent App. 18/115,596
 19. T. Gokmen, V. Kalantzis, P. Ram, C. W. Wu, K. L. Clarkson, L. Horesh, and S. Ubaru. Vsa efficient lookup, 2024
 20. X. Qu, P. Ram, H. C. Samulowitz, and U. Khurana. Foundational models for tabular data, 2024
 21. T. Salonidis, Y. Zhou, M. Abbas, P. Ram, N. B. Angel, H. C. Samulowitz, and T. Chen. Combating byzantine attacks in bilevel federated learning systems, 2024
 22. T. Parnell, M. Lazuka, A. Anghel, and P. Ram. Configuration of services across multiple cloud providers, 2024
 23. M. Abbas, Y. Zhou, P. Ram, N. B. Angel, H. C. Samulowitz, T. Salonidis, and T. Chen. Enhancing in-context learning with foundation models via few-shot linear probe calibration, 2024
 24. C. Park, P. Ram, K. L. Clarkson, and L. Horesh. Gpu-accelerated k-means clustering with transformers, 2025

Preprints

1. M. Hirzel, K. Kate, A. Shinnar, S. Roy, and P. Ram. Type-driven automated learning with lale. *arXiv preprint arXiv:1906.03957*, 2019
2. M. Feffer, M. Hirzel, S. C. Hoffman, K. Kate, P. Ram, and A. Shinnar. Navigating ensemble configurations for algorithmic fairness. *arXiv preprint arXiv:2210.05594*, 2022
3. T. Klinger, L. Liu, S. Dan, M. Crouse, P. Ram, and A. Gray. Compositional program generation for systematic generalization. *arXiv preprint arXiv:2309.16467*, 2023
4. T. Ito, L. Cocchi, T. Klinger, P. Ram, M. Campbell, and L. Hearne. The importance of positional encoding initialization in transformers for relational reasoning. *arXiv preprint arXiv:2406.08272*, 2024
5. M. I. I. Alam, P. Ram, S. Dan, H. Samulowitz, and K. Kar. On the utility of domain-adjacent fine-tuned model ensembles for few-shot problems. *arXiv preprint arXiv:2406.13720*, 2024
6. H. Fernando, H. Shen, P. Ram, Y. Zhou, H. Samulowitz, N. Baracaldo, and T. Chen. Mitigating forgetting in LLM supervised fine-tuning and preference learning. *arXiv preprint arXiv:2410.15483*, 2024
7. T. Ito, M. Campbell, L. Horesh, T. Klinger, and P. Ram. Quantifying artificial intelligence through algebraic generalization. *arXiv preprint arXiv:2411.05943*, 2024
8. I. Kang, P. Ram, Y. Zhou, H. Samulowitz, and O. Seneviratne. On learning representations for tabular data distillation. *arXiv preprint arXiv:2501.13905*, 2025

Thesis

1. P. Ram. New paradigms for approximate nearest-neighbor search. 2013